

# Domain Adaptation of Self-RAG for Vietnamese Legal QA

Trịnh Nhật Tân

Trường Đại học Công nghệ Thông tin

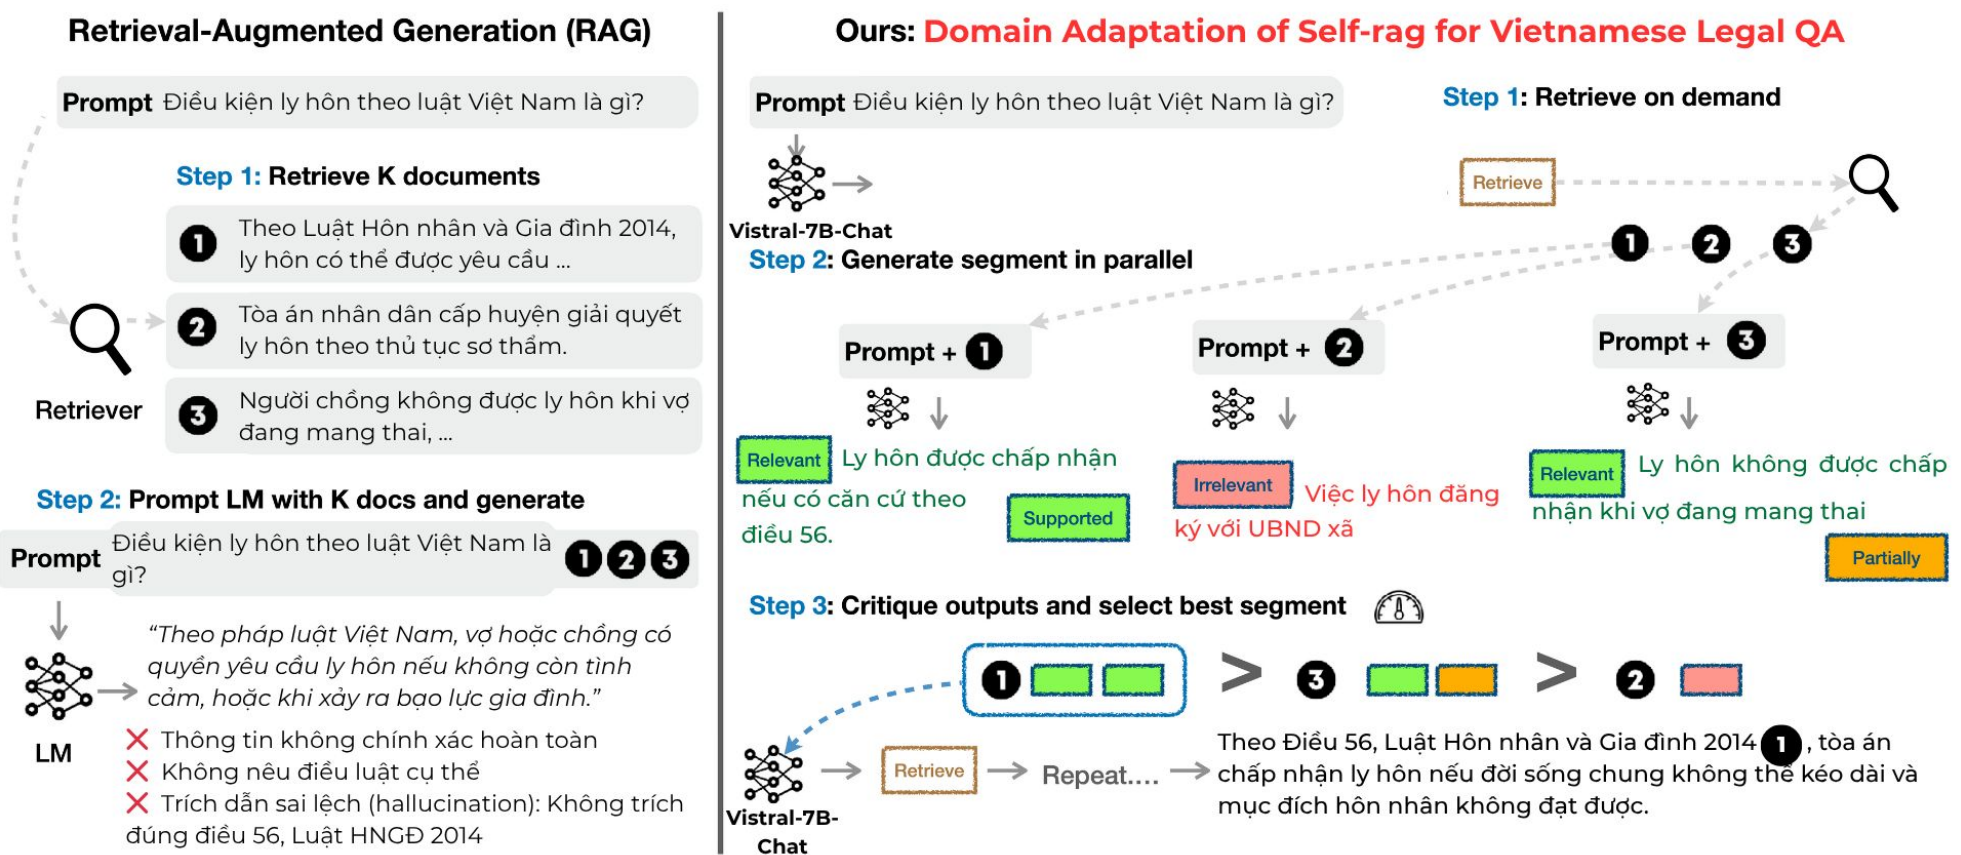
## What ?

- Apply **Self-RAG** framework to Vietnamese legal question answering.
- Use **Vistral-7B-Chat** **fine-tuned** with reflection tokens (**[Retrieve]**, **[IsRel]**, **[IsSup]**).
- Build a **legal text corpus** from TVPL, CP\_VLC, Legal Zalo, and Thuvienphapluat.vn.
- Train the model to **self-evaluate** and **retrieve only when needed**.

## Why ?

- Legal domain requires **high factual** accuracy and **citation transparency**.
- Traditional RAG often leads to **hallucinated** or **irrelevant answers**.
- Self-RAG allows **on-demand retrieval** and **self-reflection**, improving trust.
- Tailoring for Vietnamese helps address **low-resource language challenges**.

## Overview



## Description

### 1. Legal Corpus Construction

- Collected from TVPL, CP\_VLC, Legal Zalo 2021, and Thuvienphapluat.vn.
- Texts are preprocessed: normalized, segmented, and enriched with metadata (date, legal domain, issuer).
- Passages are split (~100 tokens) for retrieval and QA training.

### 2. Reflection Token Training

- QA pairs are generated using GPT-4, based on real legal passages.
- GPT-4 is prompted to annotate reflection tokens (**[Retrieve]**, **[IsRel]**, **[IsSup]**) for each instance.
- A manually verified subset ensures annotation quality.
- Vistral-7B-Chat is fine-tuned (via LoRA) to learn selective retrieval and self-assessment through these tokens.

### 3. Evaluation

- Model is evaluated on a legal QA benchmark for both retrieval and answer accuracy.
- Metrics focus on citation correctness, relevance, and hallucination reduction.
- Demonstrates improved trust and transparency in domain-specific question answering.